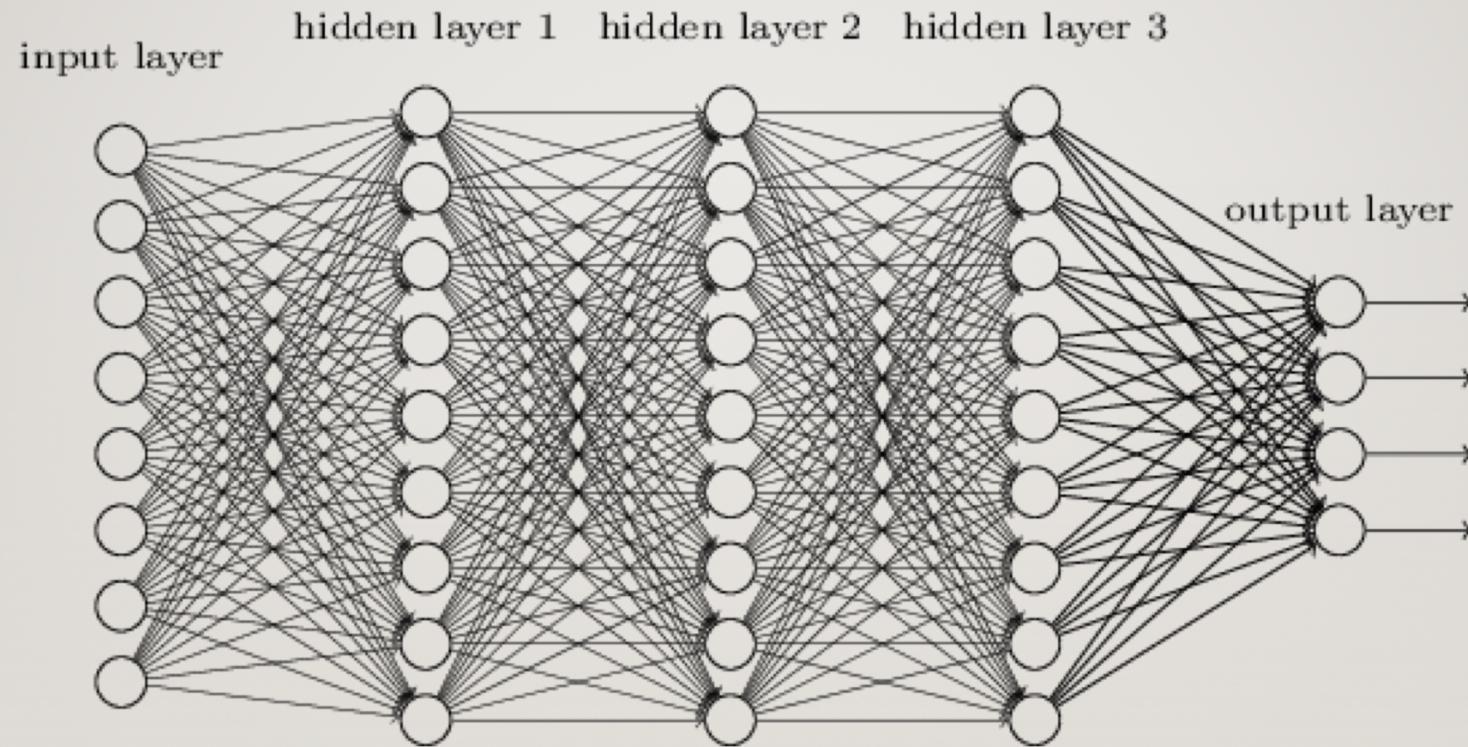


# **CHAPTER 6**

---

# **DEEP LEARNING**

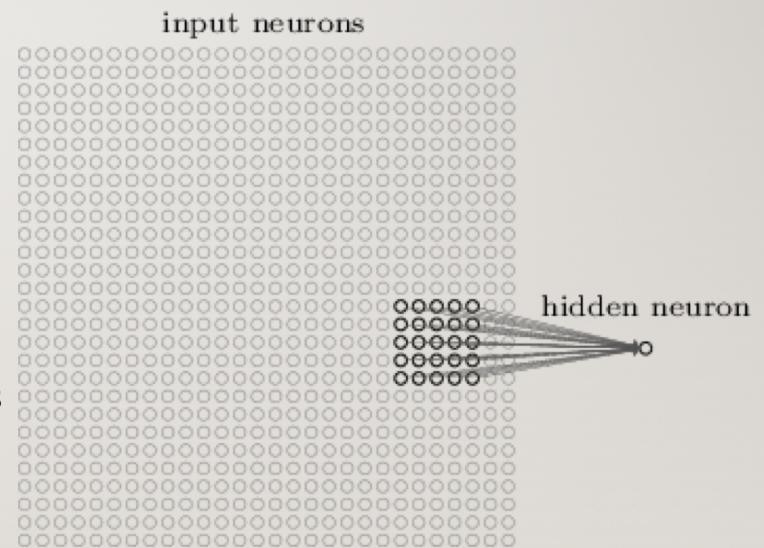


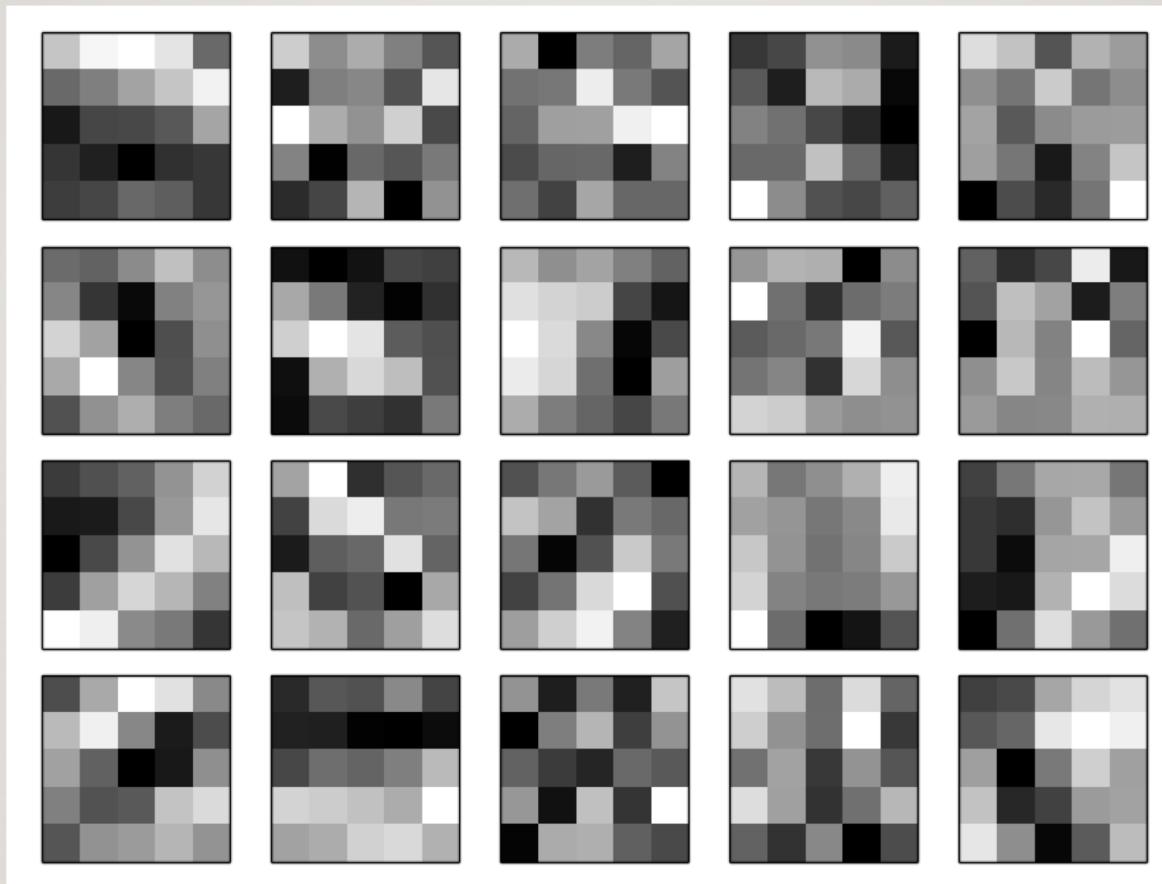


# CNNs

---

- Convolutional Neural Nets
- Spatial understanding
- Local receptive fields (filter, kernel)
- Translation invariance
- Less parameters than MLP
  - 20 5x5 kernels -> total of  $20 \times (5 \times 5 + 1) = 520$  parameters defining the convolutional layer
  - 784 input neurons, 30 hidden neurons
    - total of  $784 \times 30 + 30$  biases, for a total of 23,550 parameters

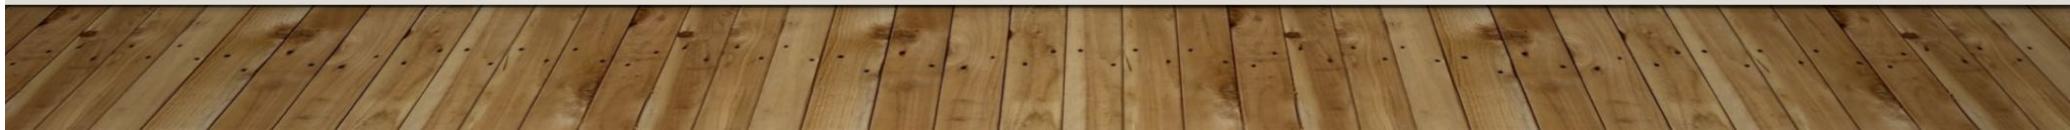
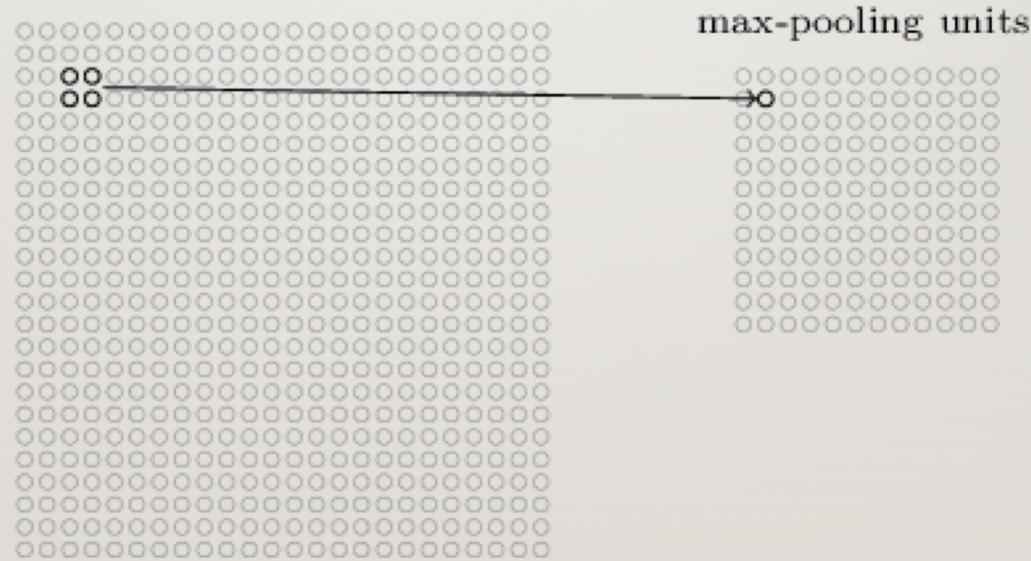




# POOLING

---

hidden neurons (output from feature map)



# MNIST

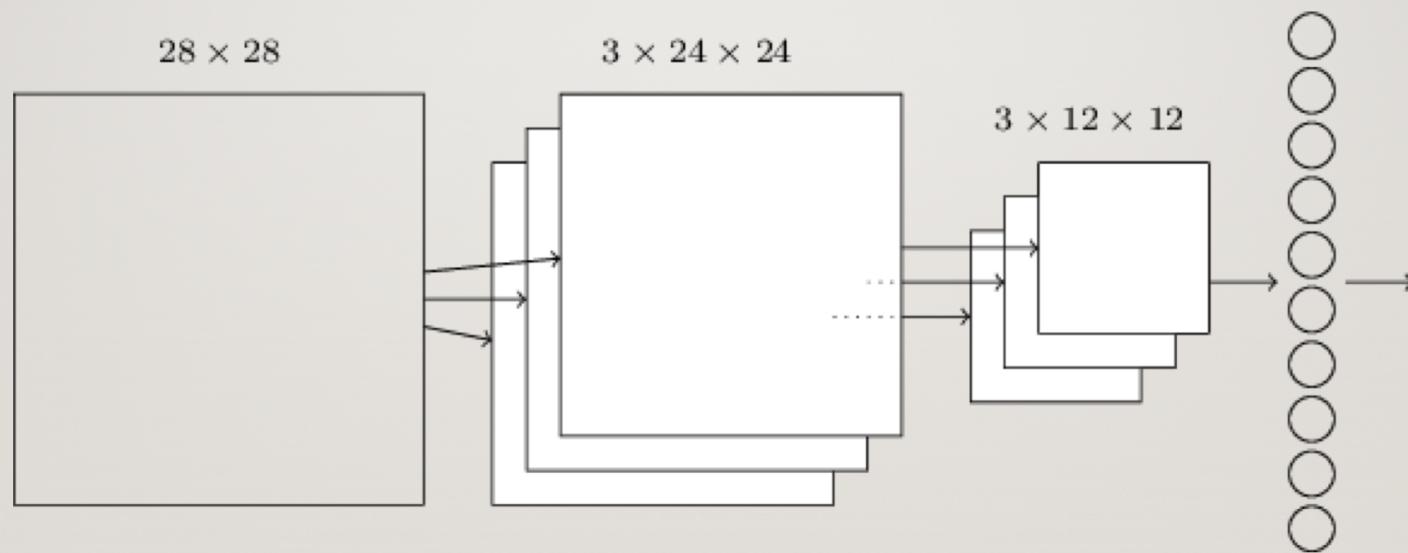
---

504 / 92



# FORWARD PROP

---



## RESULTS/NOTES

---

- More convolutional layers and more hidden layers yielded better accuracy
- Later convolutions are harder to interpret



## MAKING IMPROVEMENTS

---



# ACTIVATION FUNCTIONS

---

- sigmoid
  - The og
- Tanh
  - Trains faster, similar results
- ReLU
  - Higher accuracy! we are pretty much clueless as to why



# MORE DATA

---

- Expanded MNIST
  - 250,000 more images
- You can make more data with the data you have
  - Rotation, translation, skewing
  - Elastic distortion
  - GANs



# OTHER IMPROVEMENTS

---

- Dropout
  - Not required for convolutional layers, since they're resistant to overfitting
- Ensemble learning (kinda not really)
- Deep, Big, Simple Neural Nets Excel on Handwritten Digit Recognition, by Dan Claudiu Cireşan, Ueli Meier, Luca Maria Gambardella, and Jürgen Schmidhuber (2010)
  - hella big network, trained for a hella long time with a hella fat GPU



# BREAKTHROUGHS

---



# LMRD (2012)

---

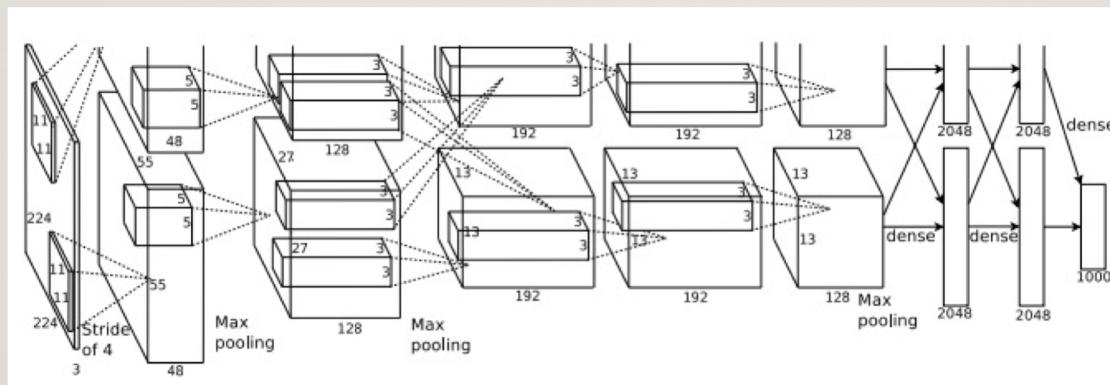
- Group of researchers from Stanford and Google
- Used a neural network to tackle ImageNet
  - 16 million+ images, 20,000+ classifications
- 9.3% to 15.8%



# KSH (2012)

---

- DCNN for a restricted subset of ImageNet in the ImageNet Large-Scale Visual Recognition Challenge (ILSVRC)
  - 84.7% for top-5, 63.3% for restrictive metric
  - Split on 2 GPUs



# KSH MODEL ARCHITECTURE

---

- Input is  $3 \times 244 \times 244$ 
  - Resize to  $256 \times 256$ , take 3 random  $244 \times 244$  crops
- First layer: 96  $11 \times 11$  kernels, stride of 4,  $3 \times 3$  max pooling with stride 2
- Second layer: 256  $5 \times 5$  kernels with max pooling
- Third, Fourth, Fifth: More convolutions without pooling
- Sixth, Seventh: Fully connected layers of 4096 neurons each
- Final: 1000-unit softmax layer
- Used ReLU, L2 regularization, dropout, momentum-based minibatch SGD
- Has inspired later work



## ILSVRC (2014)

---

- 93.33% accuracy from team based at Google
- The dude actually sat down and labeled data
- Better-than-human vision



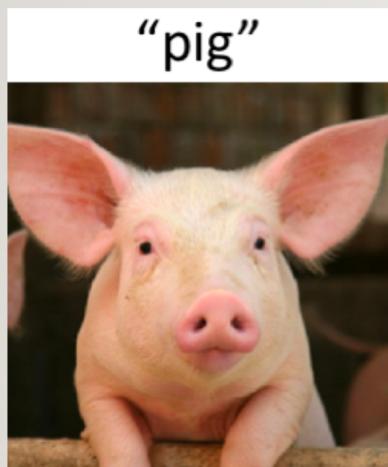
## OTHER STUFF

---

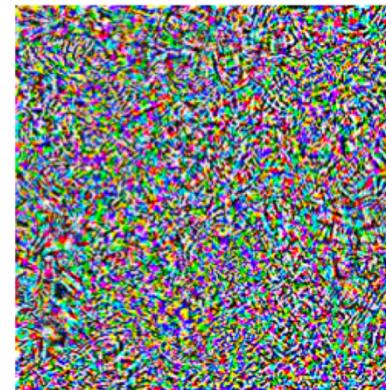


# ADVERSARIAL EXAMPLES

---



+ 0.005 x



=



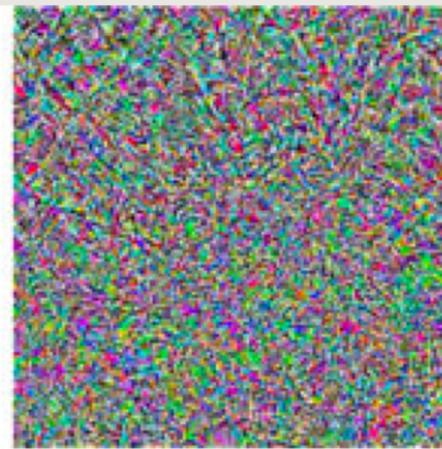
People with no idea about AI  
saying it will take over the world:

My Neural Network:





+



=



Authentic  
Input

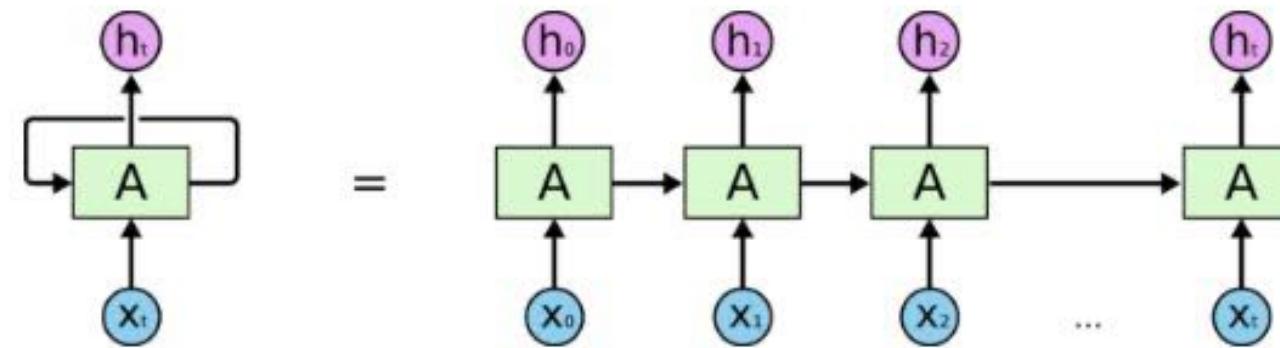
Adversarial  
Perturbation

Adversarial  
Input

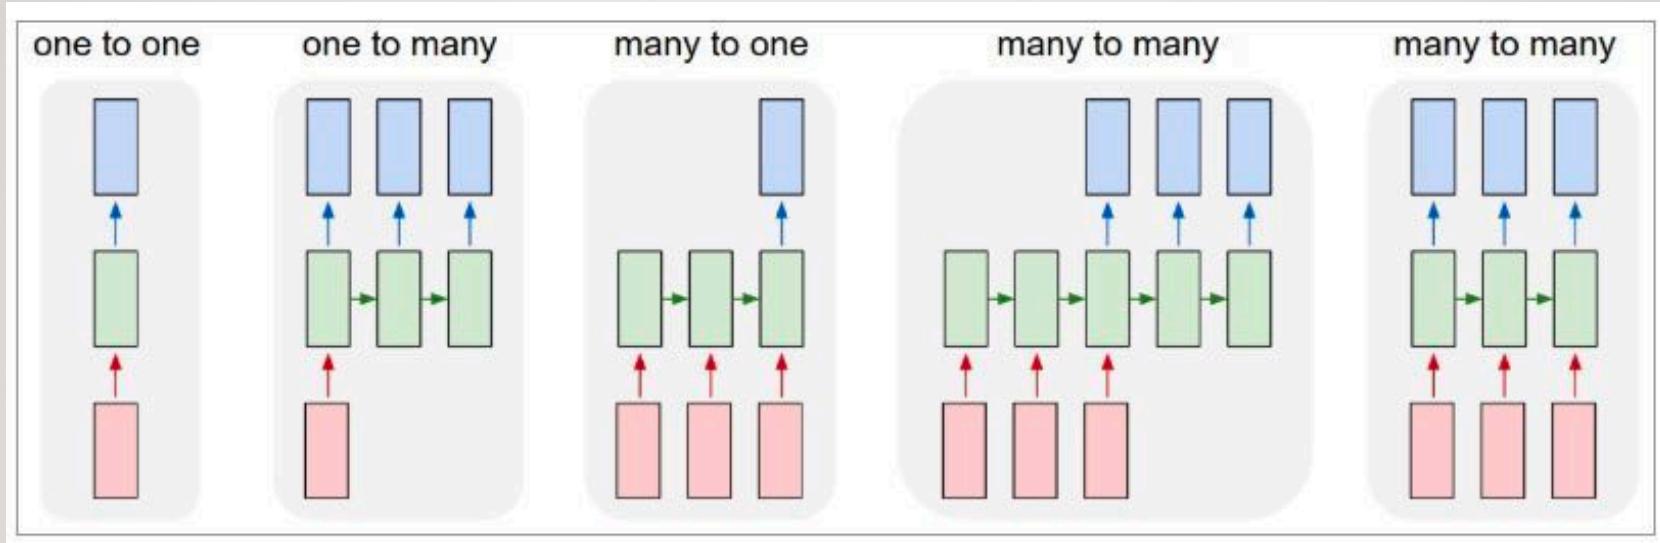
# RNNs

---

- Just like CNNs add spatial understanding, RNNs add temporal understanding



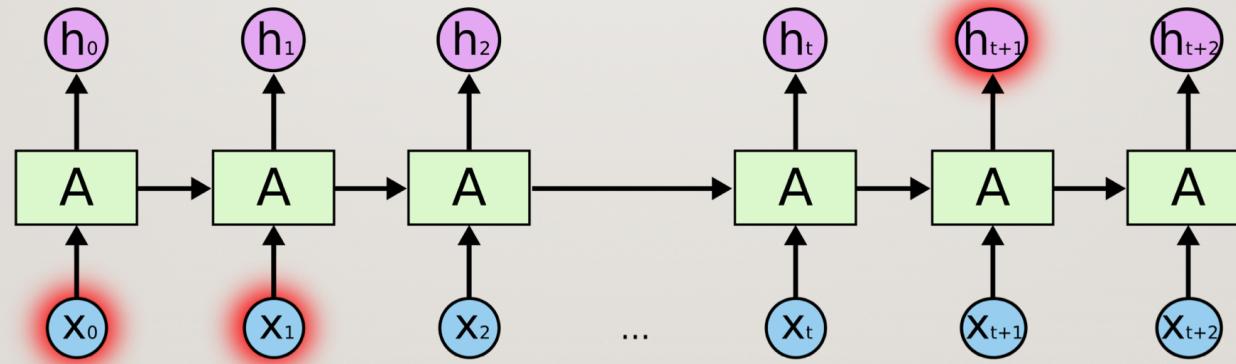
An unrolled recurrent neural network.

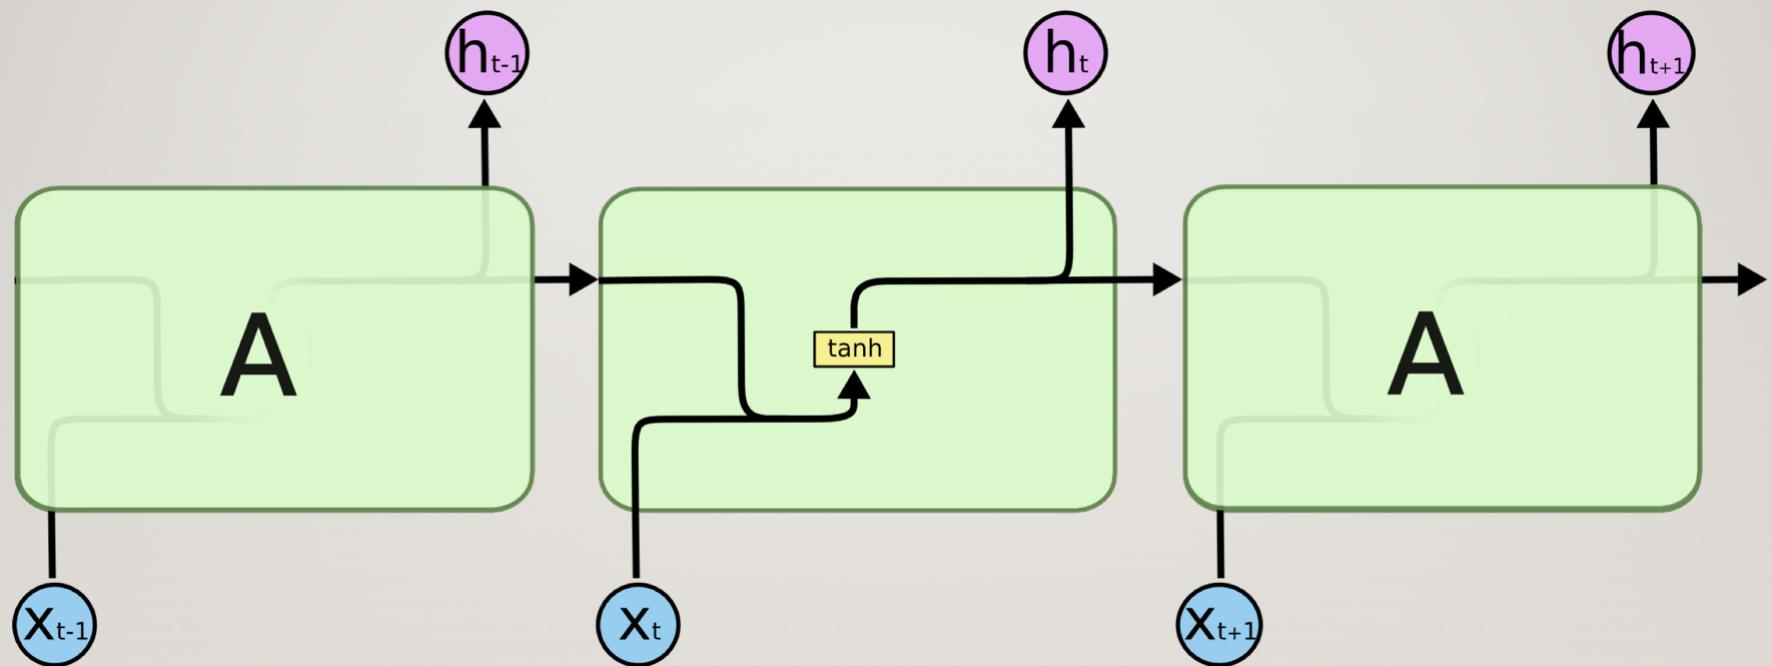


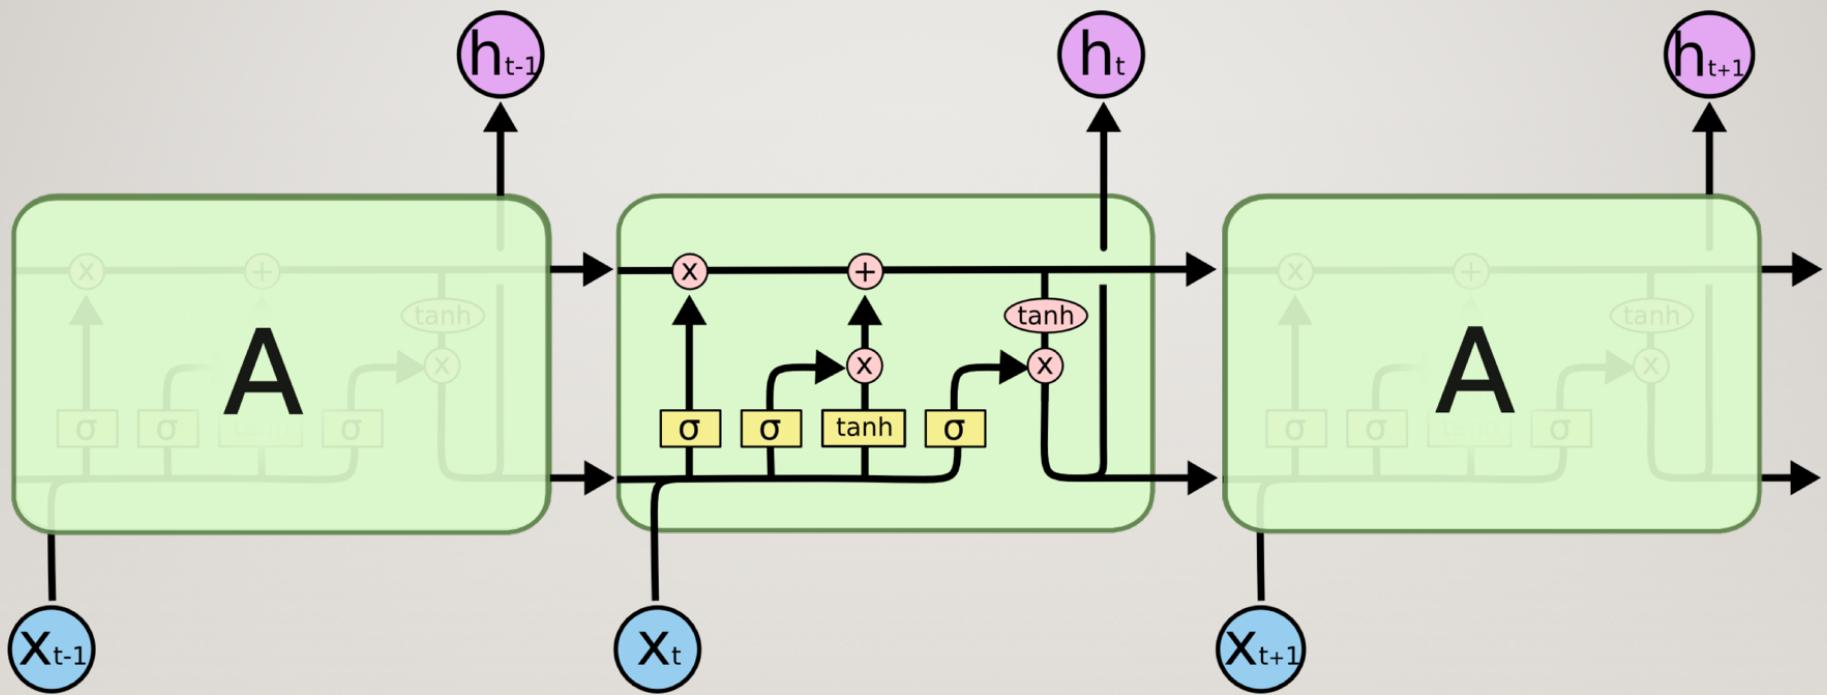
# LSTMS

---

- RNNs take soooo long to train b/c they have unstable gradients
- “I grew up in France... I speak fluent *French*.” (Long-term dependency problem)



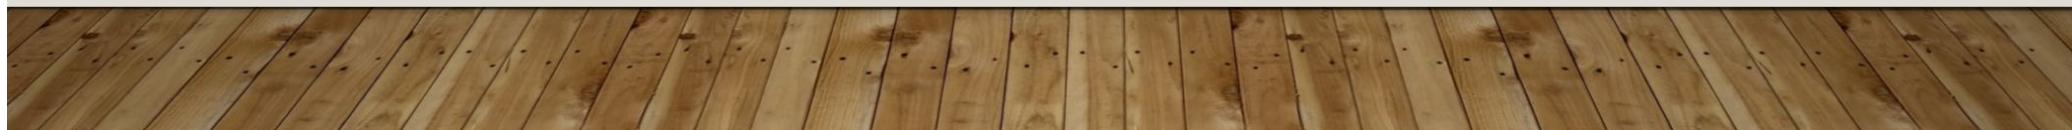




# DBNS, GENERATIVE MODELS

---

- Deep Belief Network
  - Generative model
  - Unsupervised and semi-supervised
- GANs



# INTENTION-DRIVEN USER INTERFACES

---

- Interfaces that can act on imprecision and discern the user's true intent



# FUTURE OF NN?

---

- They've done a lot of amazing things recently
- But we don't understand them nearly well enough
  - Why is it that neural networks can generalize so well?
  - How is it that they avoid overfitting as well as they do, given the very large number of parameters they learn?
  - Why is it that stochastic gradient descent works as well as it does? How well will neural networks perform as data sets are scaled?
  - Why does pooling work?



# CONWAY'S LAW

---

- “Any organization that designs a system... will inevitably produce a design whose structure is a copy of the organization's communication structure.”
- Applies to the design and engineering of systems where we have a good understanding of the different parts
- Can't be applied directly to the development of AI because we don't know what the parts are.



# FINAL THOUGHTS

---

- Medicine -> immunology, epidemiology, etc
- Deep learning is our “super-special weapon”
  - How powerful is it?
  - What other powerful idea will be needed for strong AI?
- We don’t see a lot of specialized subfields yet, everything’s built off the same ideas

