

Intro to Deep Q Networks


Jake Grigsby

University of Virginia

April 26, 2020

Agents and the Environment

Many problems in AI can be thought of as some autonomous agent interacting with an environment:



This concept is formalized as a **Markov Decision Process**...

Markov Decision Process

Definition

A Markov Decision Process (**MDP**) consists of:



- \mathcal{S} , a set of states
- \mathcal{A} , a set of actions
- $\mathcal{R} \subseteq \mathbb{R}$, a set of rewards
- a dynamics function $p : \mathcal{S} \times \mathcal{R} \times \mathcal{S} \times \mathcal{A} \rightarrow [0, 1]$

$$p(s', r | s, a) = \Pr\{\mathcal{S}_t = s', R_t = r | \mathcal{S}_{t-1} = s, A_{t-1} = a\}$$


It's common to break the dynamics function p up into a **Transition Function** $T(s, a, s') = \sum_{r \in \mathcal{R}} p(s', r | s, a)$, and a **Reward Function**

$$R(s, a) = \sum_{r \in \mathcal{R}} r \sum_{s' \in \mathcal{S}} p(s', r | s, a)$$

A Quick Example



(s,a) is called a q-state



s is a state



(s,a,s') called a transition

$$T(s,a,s') = P(s'|s,a)$$

$$R(s,a,s')$$



The RL Problem

The goal of RL agents is to find a **policy**¹ $\pi^* : \mathcal{S} \rightarrow \mathcal{A}$ that maximizes the *expected discounted return*

$$\pi^* = \underset{\pi}{\operatorname{argmax}} \mathbb{E}_{\tau \sim \pi} \left[\sum_{t=0}^{t=\infty} \gamma^t R_t \right]$$

where $\gamma \in [0, 1)$ is the *discount factor* that lets us deal with non-episodic tasks and τ is a *trajectory* (a sequence of states and actions that describe the agent's experience)

¹Policies can also be stochastic, in which case they're written $\pi(a|s) : \mathcal{S} \times \mathcal{A} \rightarrow [0, 1]$


Simplifying Assumptions

We begin by making some assumptions about the task we are trying to solve:

- 1 The dynamics of the model ($p(s', r|s, a)$) are known
- 2 $|\mathcal{S}| \ll \infty$
- 3 $|\mathcal{A}| \ll \infty$

Generalized Policy Iteration

Solution: Policy Iteration Dynamic Programming



We'll skip these details because knowledge of dynamics is such a limiting assumption in our case. More info can be found in [4]

Simplifying Assumptions

- 1 ~~The dynamics of the model ($p(s', r|s, a)$) are known~~
- 2 $|\mathcal{S}| \ll \infty$
- 3 $|\mathcal{A}| \ll \infty$

What if the environment dynamics are unknown?

Value Methods

Definition

Value methods attempt to learn the optimal *Q Function*

$$Q^*(s, a) = \max_{\pi} \mathbb{E}_{\pi} \left[\sum_{k=0}^{k=\infty} \gamma^k R_{t+k+1} \mid \mathcal{S}_t = s, \mathcal{A}_t = a \right]$$

Why? Because given $Q^*(s, a)$, the optimal policy can easily be computed by

$$\pi^*(s) = \underset{a}{\operatorname{argmax}} Q^*(s, a)$$

Value Methods: Temporal Difference

- Randomly initialize $Q(s, a)$ and use interactions with the environment as a sample to update this 'bootstrap'
- Updates based on the Bellman Equation:

$$Q^{\pi}(s, a) = \mathbb{E}_{s'} \left[r(s, a) + \gamma \mathbb{E}_{a' \sim \pi} [Q^{\pi}(s', a')] \right]$$

Value Methods: Temporal Difference

- Randomly initialize $Q(s, a)$ and use interactions with the environment as a sample to update this 'bootstrap'
- Updates based on the Bellman Equation:

$$Q^{\pi}(s, a) = \mathbb{E}_{s'} \left[r(s, a) + \gamma \mathbb{E}_{a' \sim \pi} [Q^{\pi}(s', a')] \right]$$

- We minimize the *Bellman Error*:

$$(r(s, a) + \gamma \max_{a'} Q(s', a')) - Q(s, a)$$

Value Methods: Temporal Difference

Q-learning (off-policy TD control) for estimating $\pi \approx \pi_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:

 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Take action A , observe R, S'

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$

$S \leftarrow S'$

 until S is terminal

Figure: Q-Learning Pseudo-code [4]

A Quick Note on Exploration vs. Exploitation

- At each time step, the agent must choose between "exploiting" the action it currently thinks has the best return and "exploring" alternatives to learn more about them.
- Most convergence guarantees assume state *coverage*
 - ▶ Every state will be visited an infinite number of times in an infinite number of timesteps.
 - ▶ This can be achieved by enforcing:

$$\pi(a|s) > 0, \forall s \in \mathcal{S}$$

A Quick Note on Exploration vs. Exploitation

- The simplest way to do this is to make an existing policy ϵ -greedy:

$$\pi'(s) = \begin{cases} \pi(s) & \text{with probability } (1 - \epsilon); \\ \pi_{\text{random}}(s) & \text{with probability } \epsilon; \end{cases}$$

A Quick Note on Exploration vs. Exploitation

- The simplest way to do this is to make an existing policy ϵ -greedy:

$$\pi'(s) = \begin{cases} \pi(s) & \text{with probability } (1 - \epsilon); \\ \pi_{\text{random}}(s) & \text{with probability } \epsilon; \end{cases}$$

- This can be thought of as injecting noise into the action space
- All of the agent's we'll be talking about use this general approach, but there is a lot of interesting work on motivating agents to explore efficiently. [3] [5]

Simplifying Assumptions

- 1 ~~The dynamics of the model ($p(s', r|s, a)$) are known~~
- 2 ~~$|\mathcal{S}| \ll \infty$~~
- 3 ~~$|\mathcal{A}| \ll \infty$~~

What if the state space is too large for dynamic programming?

Tasks with Large State Spaces

Example: Video Games

- Pixel input makes $|S| = \mathbb{Z}_{256}^{H \times W \times C}$
- Atari 2600 games make up one of the most popular benchmarks in modern RL.


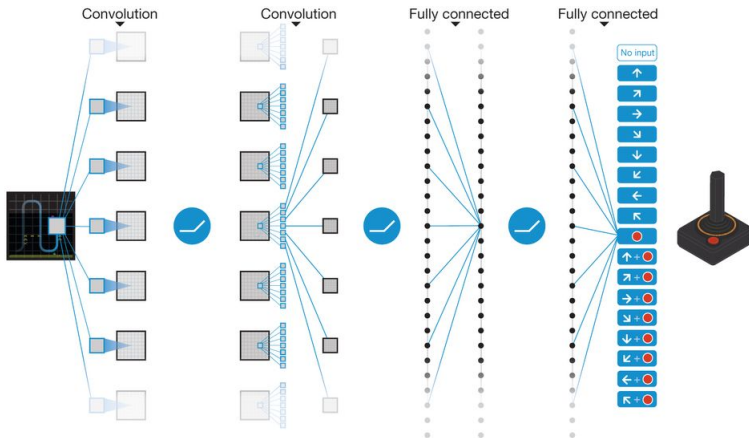


Figure: Games in the Arcade Learning Environment [1] benchmark

Deep Q Networks (DQN)

- Paramaterize Q with a neural network that can learn to recognize patterns between similar states.



Deep Q Networks (DQN)

- Train this network to minimize the *Mean Squared Bellman Error* (MSBE)

$$BE(s, a, r, s', d) = (r + \gamma(1 - d) \max_{a'} Q_{\theta'}(s', a')) - Q_{\theta}(s, a)$$

Deep Q Networks (DQN)

- Train this network to minimize the *Mean Squared Bellman Error* (MSBE)

$$BE(s, a, r, s', d) = (r + \gamma(1 - d) \max_{a'} Q_{\theta'}(s', a')) - Q_{\theta}(s, a)$$

- Kind of like supervised deep learning!
 - ▶ One important difference:
 - ★ The data distribution depends on the parameters (far from i.i.d)

Deep Q Networks (DQN)

DQNs [2] use a couple tricks to make this more like supervised learning:

- 1 Create a *replay buffer* \mathcal{R} to store transitions (s, a, r, s', d)
 - ▶ Randomly sample from this buffer at each training step.
- 2 Create a *target network* to generate the the bellman error targets.
 - ▶ This is a duplicate of the original network that is not trained but is updated with fresh params every ~ 10000 steps.

The original DQN was able to learn superhuman policies on many games with dense reward signals!

Deep Q Networks (DQN)

Algorithm 1: deep Q-learning with experience replay.

Initialize replay memory D to capacity N

Initialize action-value function Q with random weights θ

Initialize target action-value function \hat{Q} with weights $\theta^- = \theta$

For episode = 1, M **do**

Initialize sequence $s_1 = \{x_1\}$ and preprocessed sequence $\phi_1 = \phi(s_1)$

For $t = 1, T$ **do**

With probability ϵ select a random action a_t

otherwise select $a_t = \operatorname{argmax}_a Q(\phi(s_t), a; \theta)$

Execute action a_t in emulator and observe reward r_t and image x_{t+1}

Set $s_{t+1} = s_t, a_t, x_{t+1}$ and preprocess $\phi_{t+1} = \phi(s_{t+1})$

Store transition $(\phi_t, a_t, r_t, \phi_{t+1})$ in D

Sample random minibatch of transitions $(\phi_j, a_j, r_j, \phi_{j+1})$ from D

Set $y_j = \begin{cases} r_j & \text{if episode terminates at step } j+1 \\ r_j + \gamma \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta^-) & \text{otherwise} \end{cases}$

Perform a gradient descent step on $(y_j - Q(\phi_j, a_j; \theta))^2$ with respect to the network parameters θ


Every C steps reset $\hat{Q} = Q$

End For


End For

Figure: [2]

Results



Since 2015...








Single-actor agents

- A DQN
- B Prioritised Dueling
- C Rainbow
- D C51-IDS
- E FQF

Distributed agents

- F ApeX
- G R2D2
- H NGU
- I MuZero
- J Agent57

References I

-  Marc G Bellemare et al. “The arcade learning environment: An evaluation platform for general agents”. In: *Journal of Artificial Intelligence Research* 47 (2013), pp. 253–279.
-  Volodymyr Mnih et al. “Human-level control through deep reinforcement learning”. In: *Nature* 518.7540 (2015), pp. 529–533.
-  Yuri Burda et al. “Large-scale study of curiosity-driven learning”. In: *arXiv preprint arXiv:1808.04355* (2018).
-  Richard S Sutton and Andrew G Barto. *Reinforcement learning: An introduction*. MIT press, 2018.
-  Adrià Puigdomènech Badia et al. “Never Give Up: Learning Directed Exploration Strategies”. In: *arXiv preprint arXiv:2002.06038* (2020).